



# *DevOps: Puppet*

Stephen Mouring  
*Director of Software Engineering*

**Séquoia Holdings**

# Introductions



# Introduction

- Software Developer
- Director of Software Engineering for Sequoia Holdings Inc.
- Full Stack Development
  - Recent focus on big data and the cloud
- Current project is transitioning to DevOps model
- Utilizing Puppet for Infrastructure Configuration Management



# Introduction: Sequoia

- Employee Owned
- Broad competencies
- AWS and BigData specialty
- IC customer branching out into the commercial space
  - Amazon Partner
  - Cloudera Partner
- Currently involved in AWS transition efforts across four agencies



# Overview





# Overview

- What is DevOps?
- What is Puppet?
- Puppet Concepts
- Demo
- AWS/C2S Specifics



# What is DevOps?



# Intersection of Software **Dev**elopment and System **Op**erations





*Methodology*

Not A

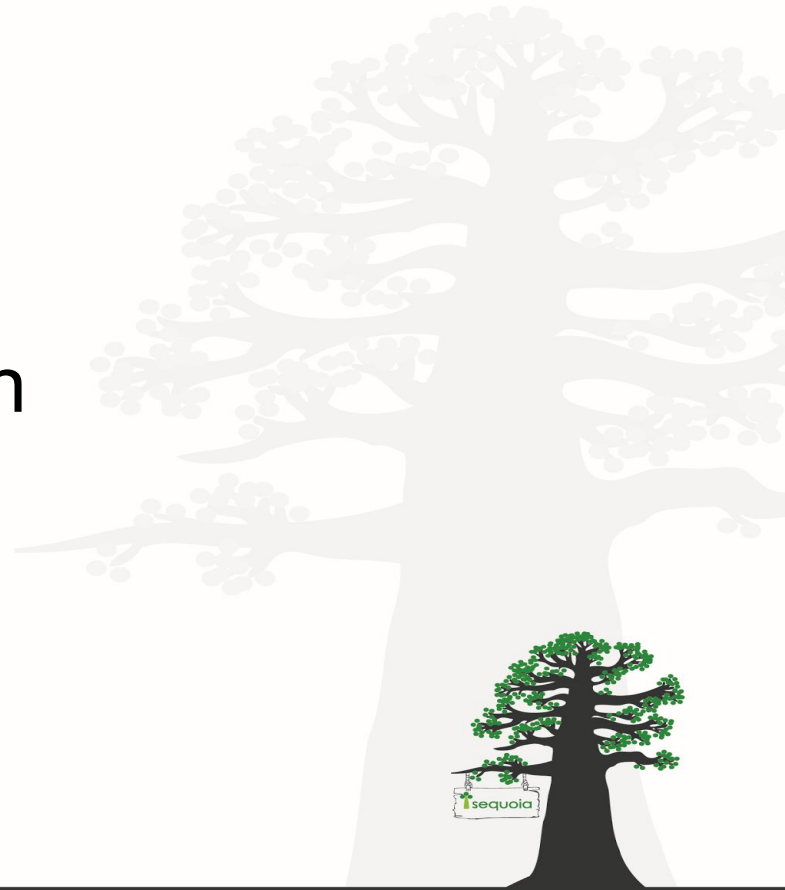
*Technology*



Leverages the best in multiple disciplines



# Automation



# Repeatability

preserving institutional knowledge



# Frameworks

that generalize specialized knowledge



# What is Puppet?



# Server configuration **automation**





# Server configuration **enforcement**



User specifies desired server state...



Puppet ensures state is applied...



... and remains applied



# Example

Raspberry Pi configuration



Why?



# Manual steps do not **scale**

especially for recurring tasks like patching!





Manual steps are **error prone**



# Automation enables **new behavior**

such as AWS autoscaling



*Declarative*  
Not  
*Imperative*



(Which means you tell it **what** *not how...*)



State specified using Puppet language



Can either run standalone or client/server



# Puppet Concepts





Suffers from disorganized documentation...



... and (deliberately?) indirect marketing



# Concept: **Resource**



# Configuration building block



```
<TYPE> { '<NAME>' :  
    <ATTRIBUTE>,  
    <ATTRIBUTE>,  
    . . .  
}
```



```
File { '/tmp/hello.txt':  
  owner => centos,  
  group => users,  
  mode  => 755,  
  content => 'Puppet was here.',  
}
```



# Concept: **Resource Type**





# Define Resource attributes and implementation



```
<TYPE> { '<NAME>' :  
    <ATTRIBUTE>,  
    <ATTRIBUTE>,  
    . . .  
}
```



Several built in:  
File, Package, Service, User, Group ...



You can also write your own



Resource Types are a lens for viewing system configuration



# DEMO:

# Examining Resource Types



# Concept: **Class**



A Class is a **group** of related **resources**





```
class <NAME> {  
    <RESOURCE>  
    <RESOURCE>  
}
```



```
class tomcat {  
  User { 'tomcat':  
    group => 'uesrs'  
  }  
  Package { 'apache-tomcat-7':  
    ensure => present  
  }  
  ...  
}
```



A Class should have a single purpose



Classes are **defined**



Then they must be **declared**



```
include <NAME>
```



```
include tomcat
```



**Defining** creates a class





**Declaring** indicates a class should be applied



Declaring a Class applies all contained resources



(at the **same time!**)

unless you specify otherwise



# Concept: **Node**



A Node represents a particular server



Node contains Resources and Classes



```
node '<NAME>' {  
    <RESOURCE>,  
    ...  
  
    <CLASS DECLARATION>,  
    ...  
}
```



```
node 'example01.example.com' {  
  User { 'foo':  
    home => '/home/foo'  
  }  
  
  include tomcat  
}
```





Node determine which configuration  
is applied to which servers



The Node name matches a server's hostname



(‘default’ and wildcards also available)



# Concept Review



# Concept: **Manifest**



A file that contains:

- Resources
- Classes (definitions / declarations)
- Nodes



Ends in .pp



(we will build on this in a moment)





# Puppet Architecture



# Standalone



# Client / Server



# (But they called it Master/Agent)

*Because Puppet Master was too good a phrase to not use...*



In Standalone mode you explicitly invoke Puppet  
with a Manifest



In Master/Agent mode you have a **site manifest**



An Agent “checks in” with the Master



The Master computes the “catalog” that applies  
to the Agent





# Concept: **Modules**



A Module is a way of distributing and organizing  
a manifest and support files



# A directory structure and convention



## <MODULE DIRECTORY>

- > *manifests*
- > *files*
- > *templates*
- > *lib*
- > *facts.d*
- > *examples*
- > *spec*



The `modules` directory must contain an `init.pp` manifest



`init.pp` must contain a class with the same  
name as the module



To use a module, you:

```
include <module_name>
```



This doesn't include the “module”





# It includes a class...

... which puppet finds using the module directories



You can readily bundle files and other resources  
in modules



# Concept: **Environments**



You can create multiple site manifests / sets of modules



When a node registers, it can be assigned an environment



# Concept Review



# DEMO:

## **Poking around file structure**



# Puppet Setup





# Standalone



# Invoke a manifest explicitly

(with an optional module path)



**Pros:**  
Simple



**Cons:**  
Hard to scale, computation on the client



# Master/Agent



# Install Puppet Master on a server



# Puppet Master has site manifest

(or several if you have multiple environments)



# Install Puppet Agent on servers





# Puppet Agent gets a node name

(and optionally an environment)



This is based on DNS by default



Each Agent must be registered with  
the Master and approved



# Master can be configured to auto approve

(enter tools like Foreman and IPA)



Puppet Agent then polls for configuration changes on Master



## **Pros:**

Computation on Puppet Master, central copy of configuration



**Cons:**  
Complex, must manage DNS



# DEMO: Puppet In Action





# AWS Considerations



# AMI Management / Autoscaling



# Security Compliance



# Control / Team Independence



# Questions?



**[smouring@sequoiainc.com](mailto:smouring@sequoiainc.com)**

